



# Mastering Observability with OS Tools



**Luca Raveri**

Software Engineer at Talentware

# Who Am I

- **Hi**👋 I'm Luca Raveri, Software Engineer and AWS Solutions Architect



💻 Backend Development

☁️ Cloud Architectures

🔍 Observability

📌 Venice, Italy



# Talentware

- **Mission:** Helping companies build, retain, and grow talent through a skill-based approach.
- **Tech Stack:** Full-stack JavaScript, AWS, and an internal Data Science team.

 Milan, Italy



**Do you  
really know  
your system?**



# Problem Description

- We run a complex software system in production.
- With limited testing, some bugs still make it to release.
- Metrics show anomalies but don't reveal the root cause.
- Logs are massive and hard to analyze.
- Troubleshooting takes days and often requires downtime.

# What are Monitoring and Observability?

- **Monitoring**: Continuous measurement of a system's health to detect anomalies, performance issues, or downtime.
- **Observability**: The ability to understand the internal state of a complex system by examining the data it produces (logs, metrics, traces).

*Key difference:*

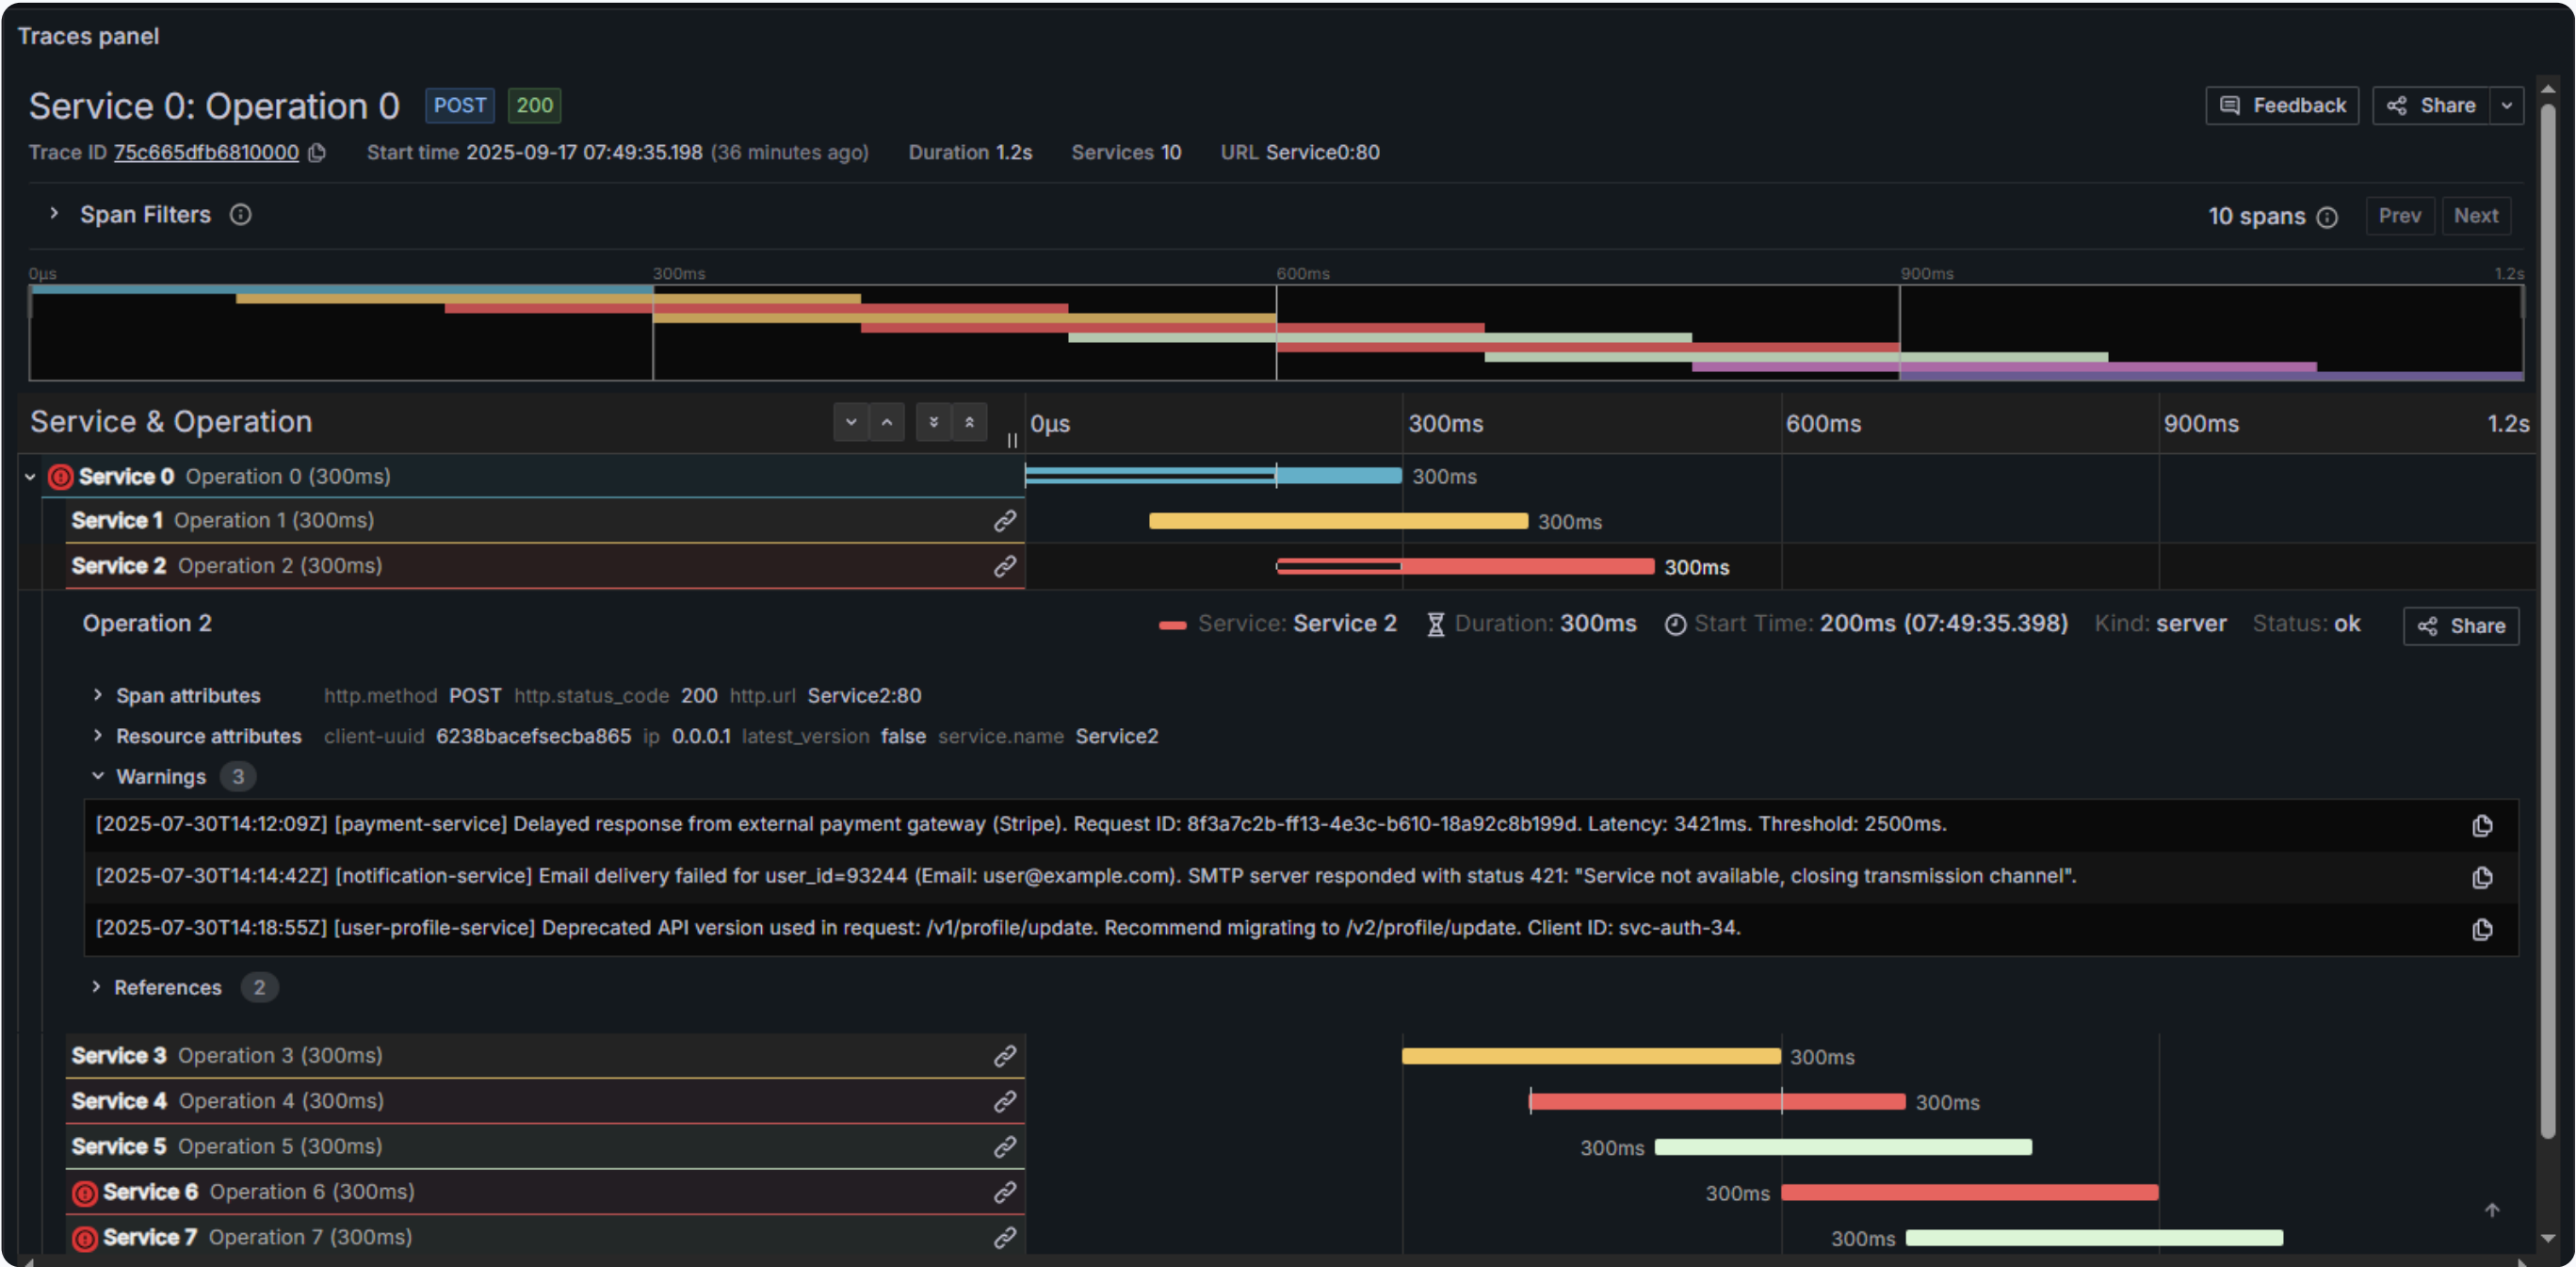
- Monitoring tells you that *something is wrong*.
- Observability helps you understand *why it is wrong*.

# The Three Pillars Of Observability

- **Metrics:** Numeric measurements over time (e.g., CPU usage, latency, error rates) → *What is happening*
- **Logs:** Discrete, timestamped records of events (e.g., errors, transactions, system messages) → *Why is happening*
- **Traces:** Detailed records that track the flow of a request through a system, showing how components interact → *Where is happening*



# What is a Trace





# Key Benefits of Observability

- **Lower MTTD (Mean Time to Detect)**

Problems are identified quickly, often before users notice.

- **Lower MTTR (Mean Time to Resolve)**

Root causes are found and fixed faster, reducing downtime and impact.

# Other Advantages

- **Dev Experience**

Faster debugging, maintainable code

- **Performance**

Understand current metrics and detect bottlenecks

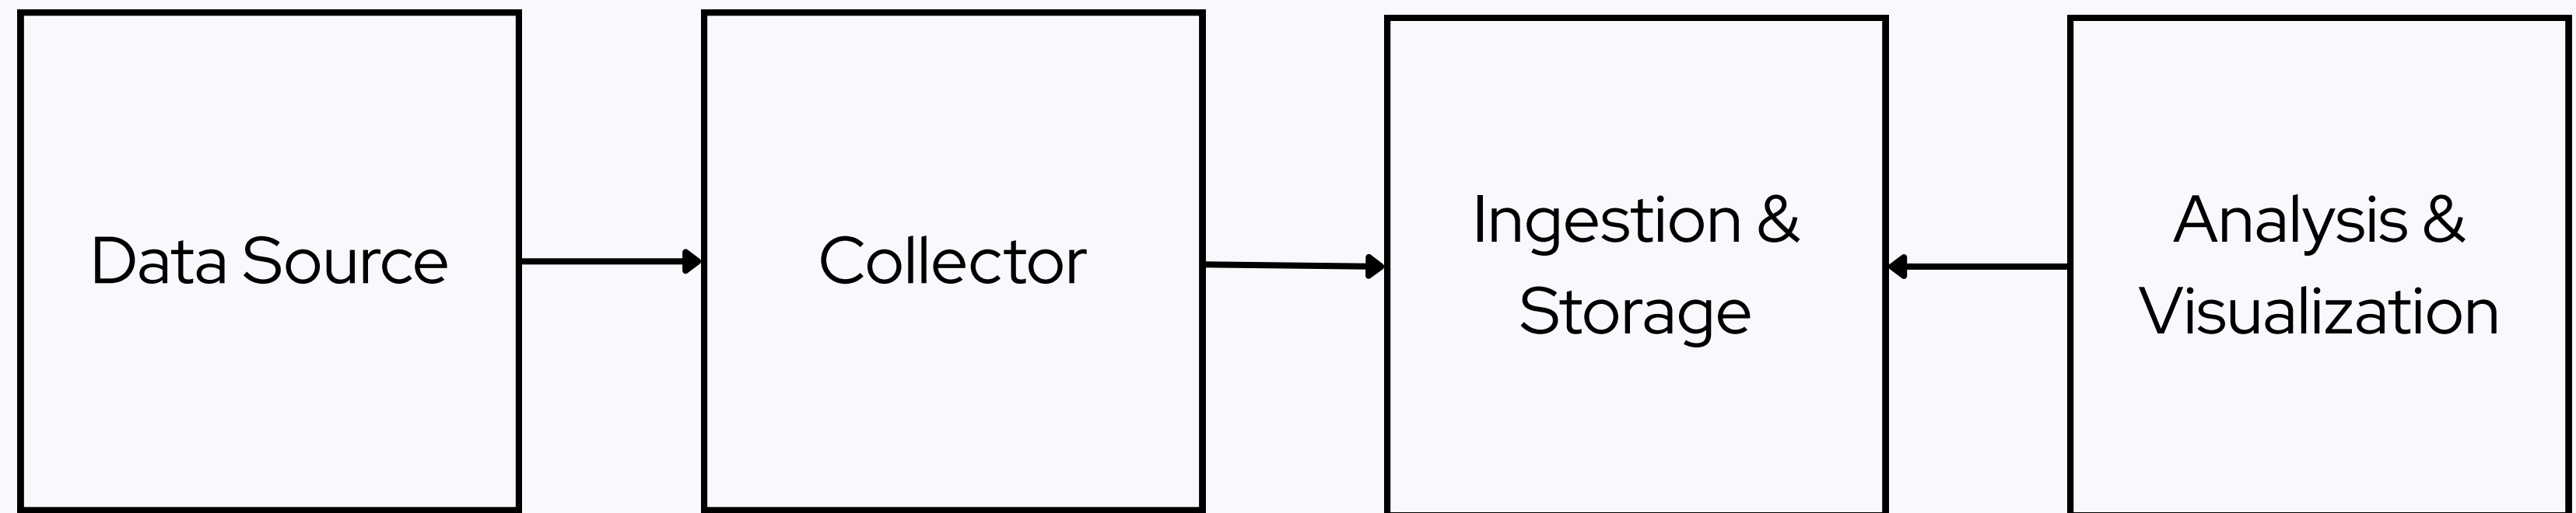
- **Business**

Control costs, track features, ensure compliance, make data-driven decisions

# Demo Time!



# Let's Build an Observability System



# LGTM Stack

## Loki



Database  
for logs

## Grafana



Visualization  
tool

## Tempo



Database  
for traces

## Mimir



Database  
for metrics

# Open Source: Pros & Cons

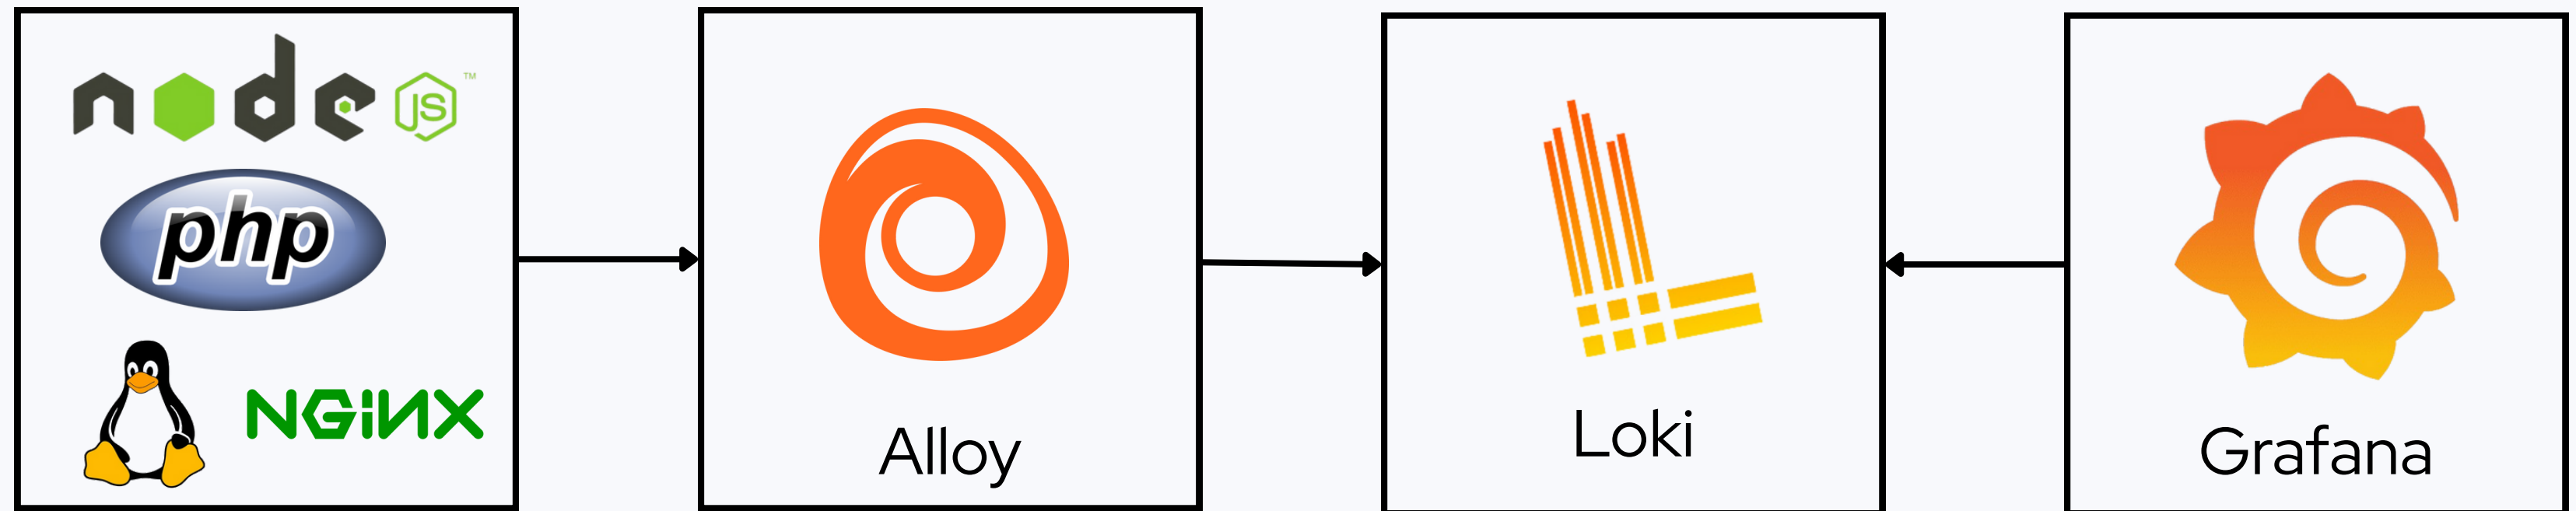
## Advantages

- No license costs
- High flexibility and customization
- Strong community support
- Avoid vendor lock-in

## Disadvantages

- Higher setup and maintenance effort
- Steeper learning curve
- Limited official support
- Integrations may require extra work

# Let's build an Observability System





# Let's Start Logging

- Use a logger instead of native functions
- Apply log levels and timestamps
- Format logs (JSON)
- Never log secrets and PII
- Set retention policies (e.g., logrotate)

# Logging Strategies

- Capture boundaries: log what enters and leaves the system
- Centralize: handle logs in a single, consistent layer if possible
- Correlate: attach a unique ID to every request/transaction
- Balance: more logs mean more insights, but also more overhead (tip: use log levels)



```
console.log(`Customer ${customerName} purchased ${itemsPurchased} items`);
```



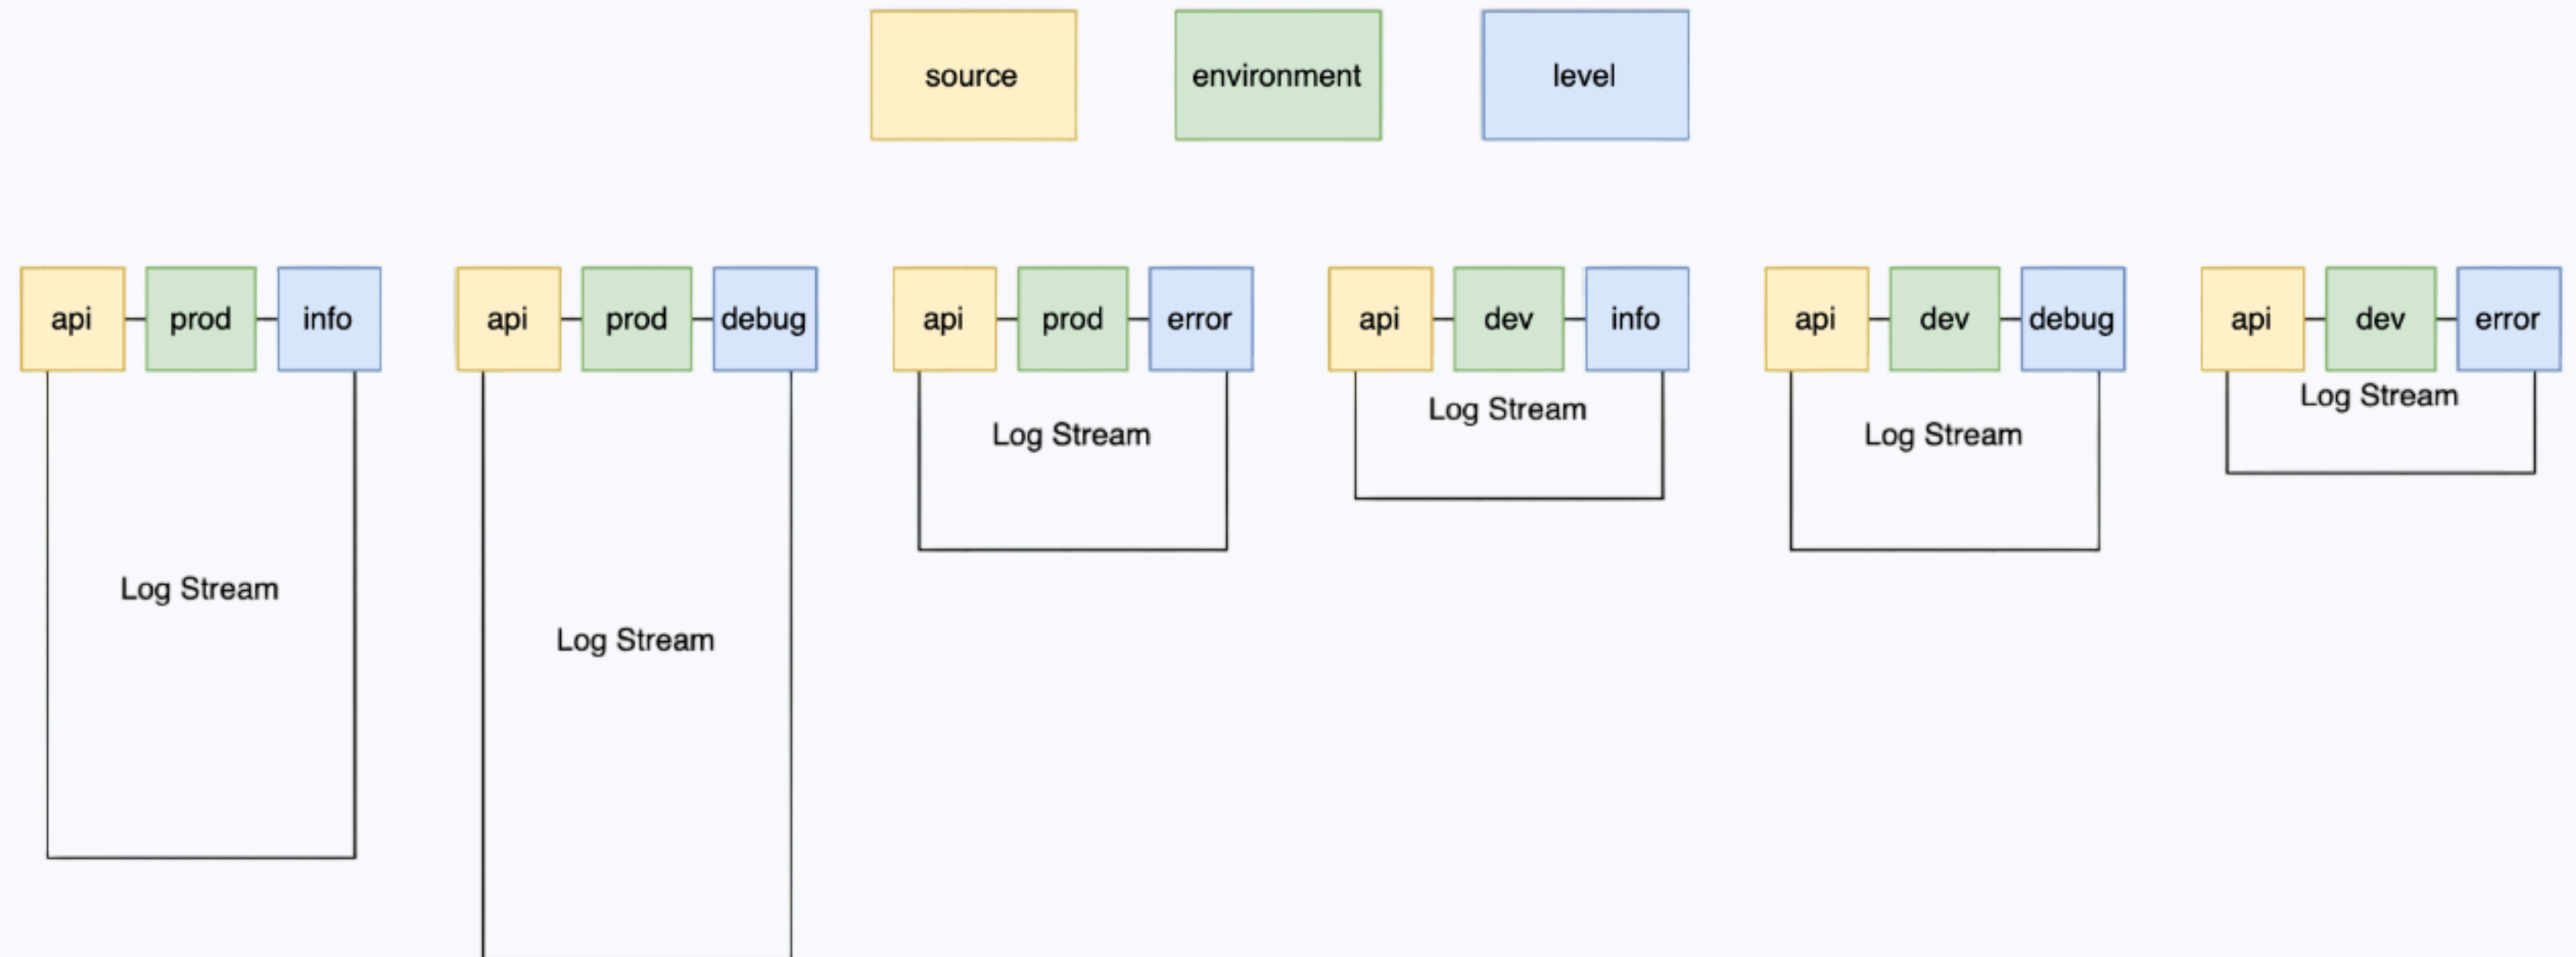
```
logger.info('Customer purchase', {  
  customerId: customer.id,  
  customerName: customer.name,  
  itemsPurchased: customer.itemsPurchased,  
  totalAmount: customer.totalAmount,  
});
```



# What is Loki

*“Loki is a horizontally scalable, highly available, multi-tenant log aggregation system inspired by Prometheus. It is designed to be very cost effective and easy to operate. It does not index the contents of the logs, but rather a set of labels for each log stream.”*

# How Loki Works



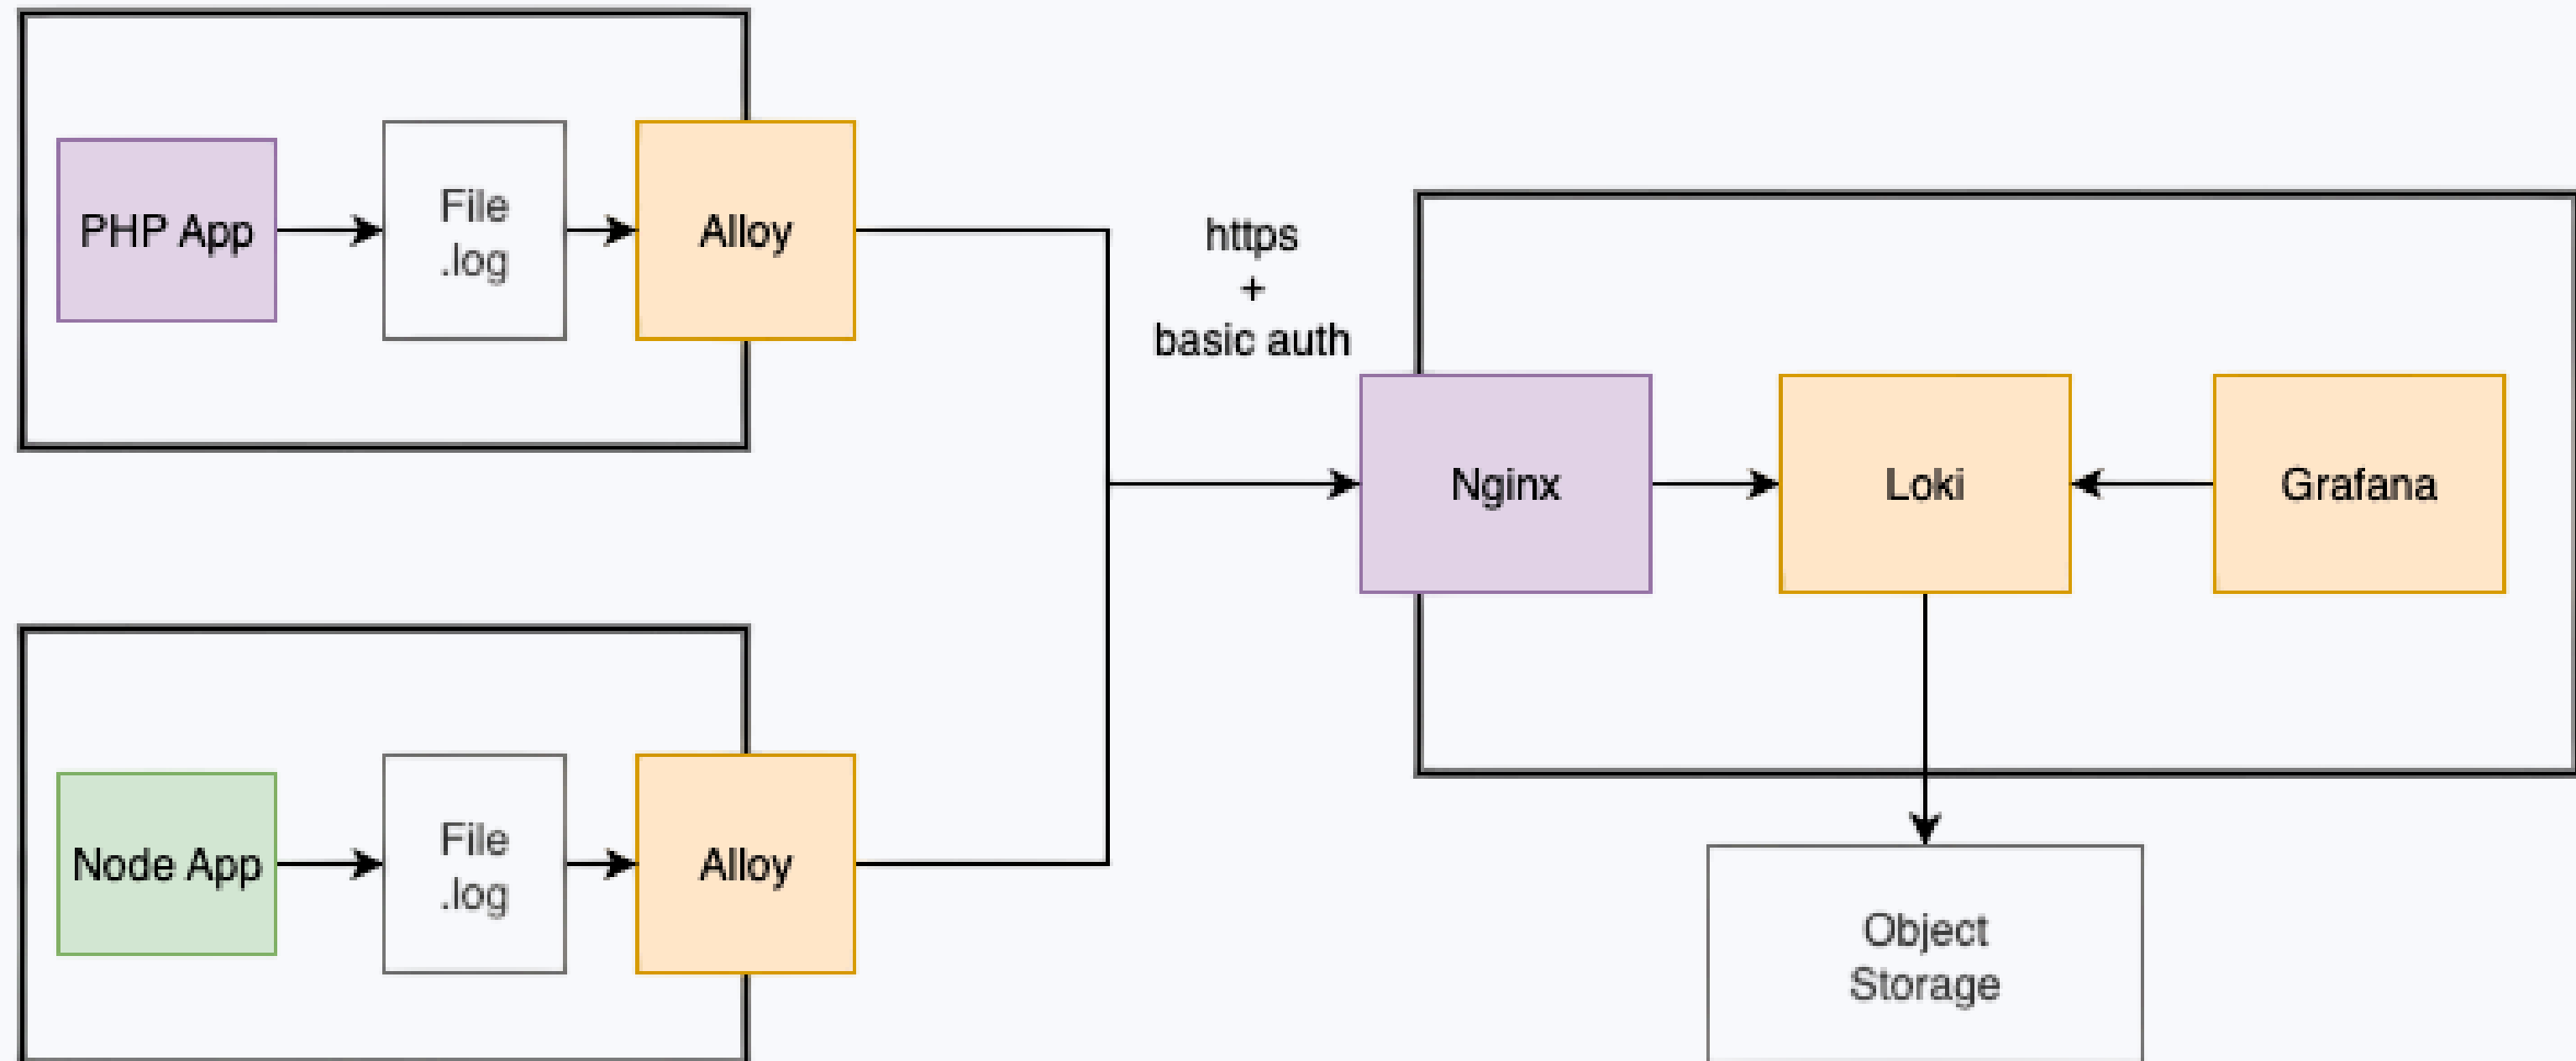
# How to Send Logs to Loki

**Grafana Alloy** is a collector that automatically reads your log files and pushes them to Loki. It can:

- Apply **static labels** to all logs from a file
- Extract **dynamic labels** from log content
- Perform **transformations** on log data



# Hosting Loki





# How to Query Logs

```
{source="grafana-demo", environment=~"${environment}", level=~"${level}"
|~ "${data}"
| json
| msg =~ ".*${message}.*"
| payload_requestId =~ ".*${requestId}.*"
```

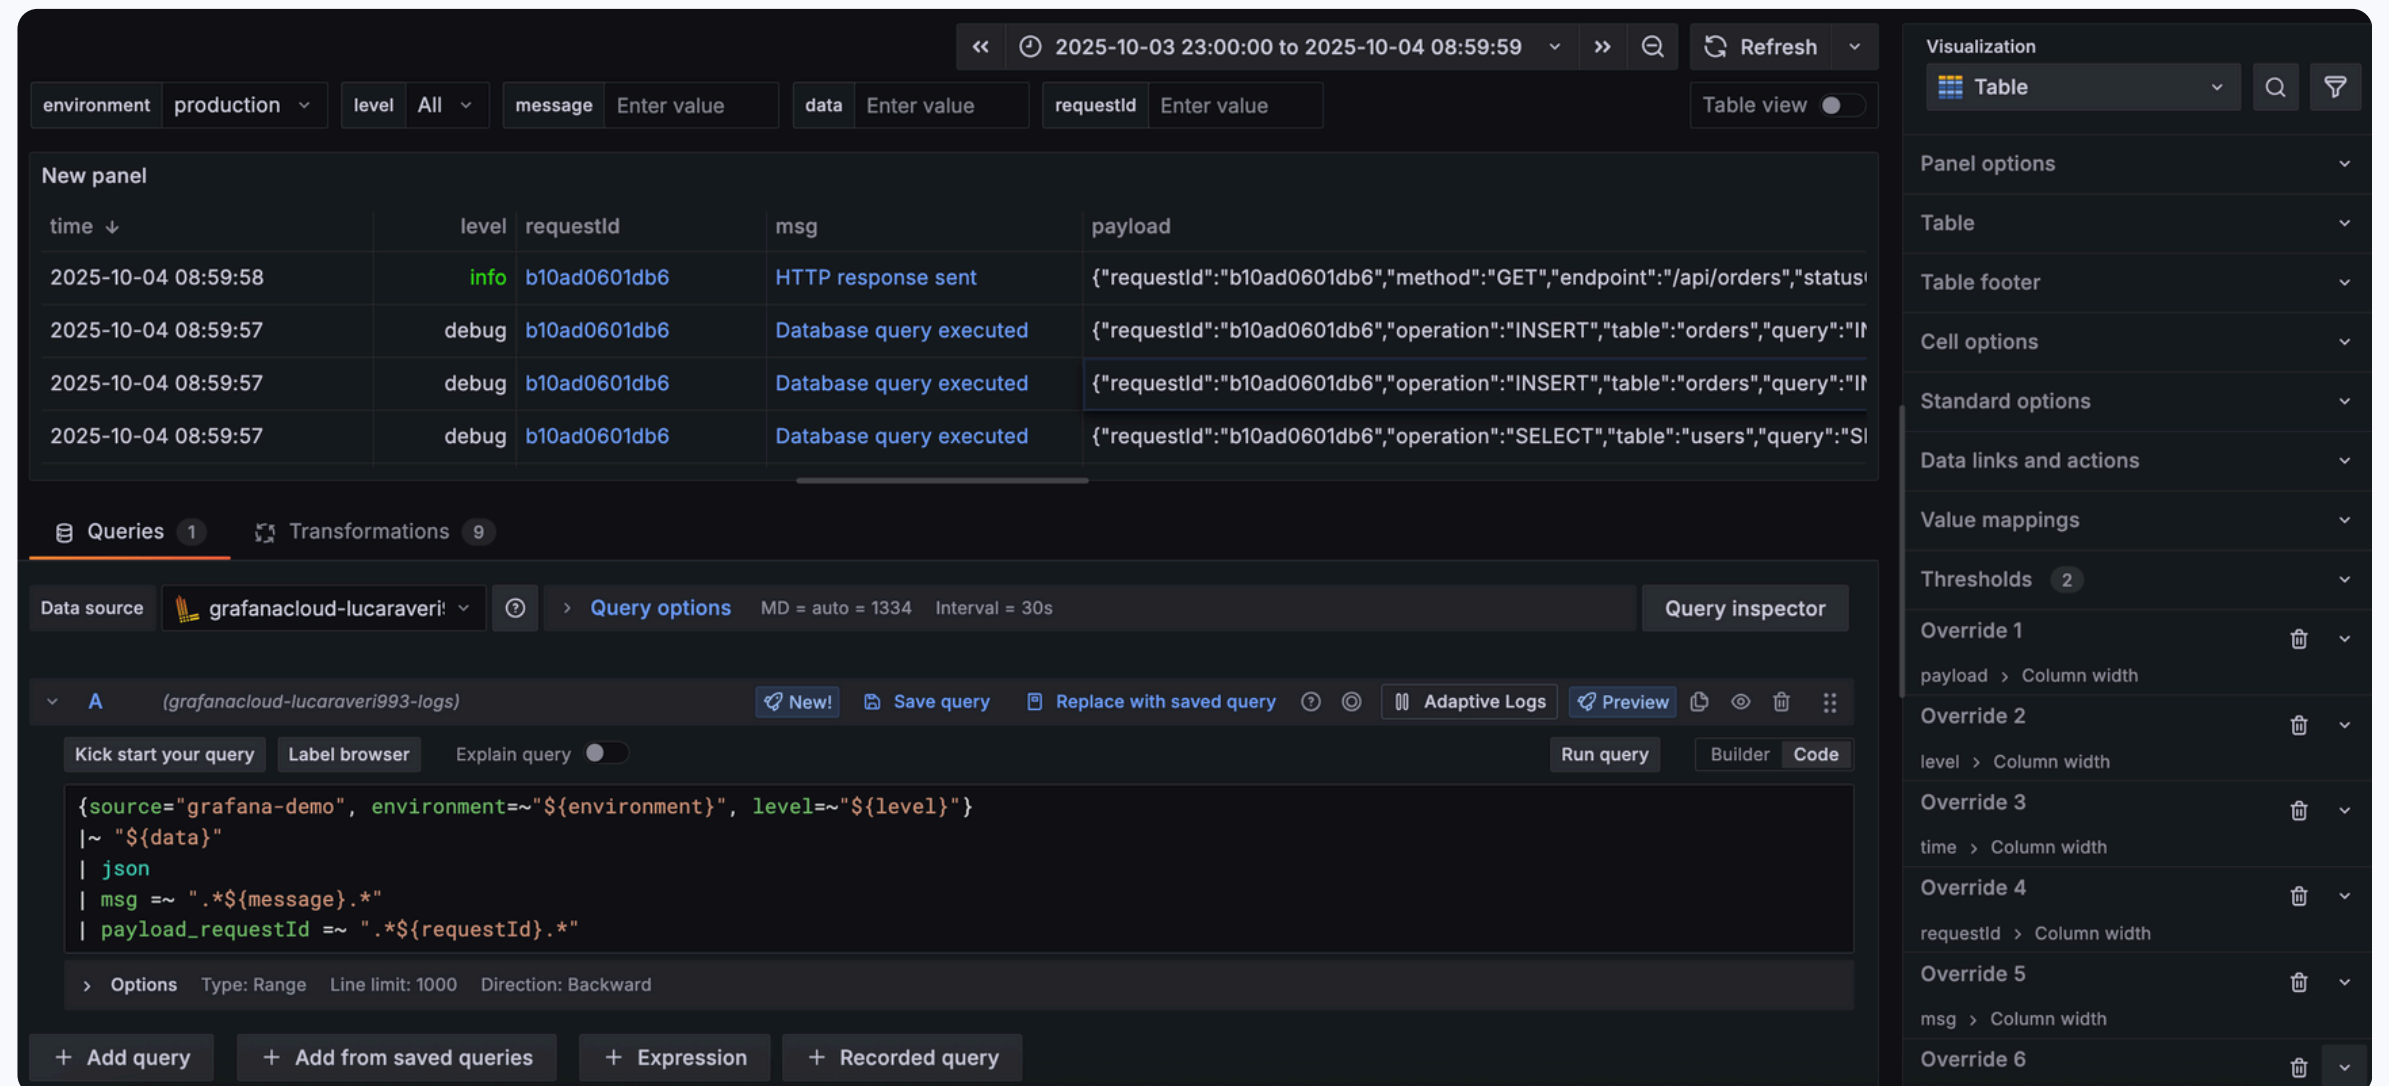


```
{job="apache_error_log", environment=~"${environment}", source=~"${source}"  
|~ "${data}"  
| regexp "\\[(?P<timestamp>[^\]]+)\]\\.[^\\[php7:(?P<level>[^\]]+)\]\\.[^PHP (?P<type>[:]+): (?P<message>.*) in (?  
P<file>[A-Z]:\\\\\\\\[^\"st]+)?(?:::(?P<line>\\\\d+)| on line (?P<line_alt>\\\\d+))]"  
| line_format "{ printf `{"timestamp": \"%s\", \"level\": \"%s\", \"type\": \"%s\", \"message\": \"%s\",  
\"file\": \"%s\", \"line\": \"%s\"}` .timestamp .level .type .message .file (or .line .line_alt) } }"  
| message =~ ".$${message}."  
| file =~ ".$${file}."  
| message != ""
```



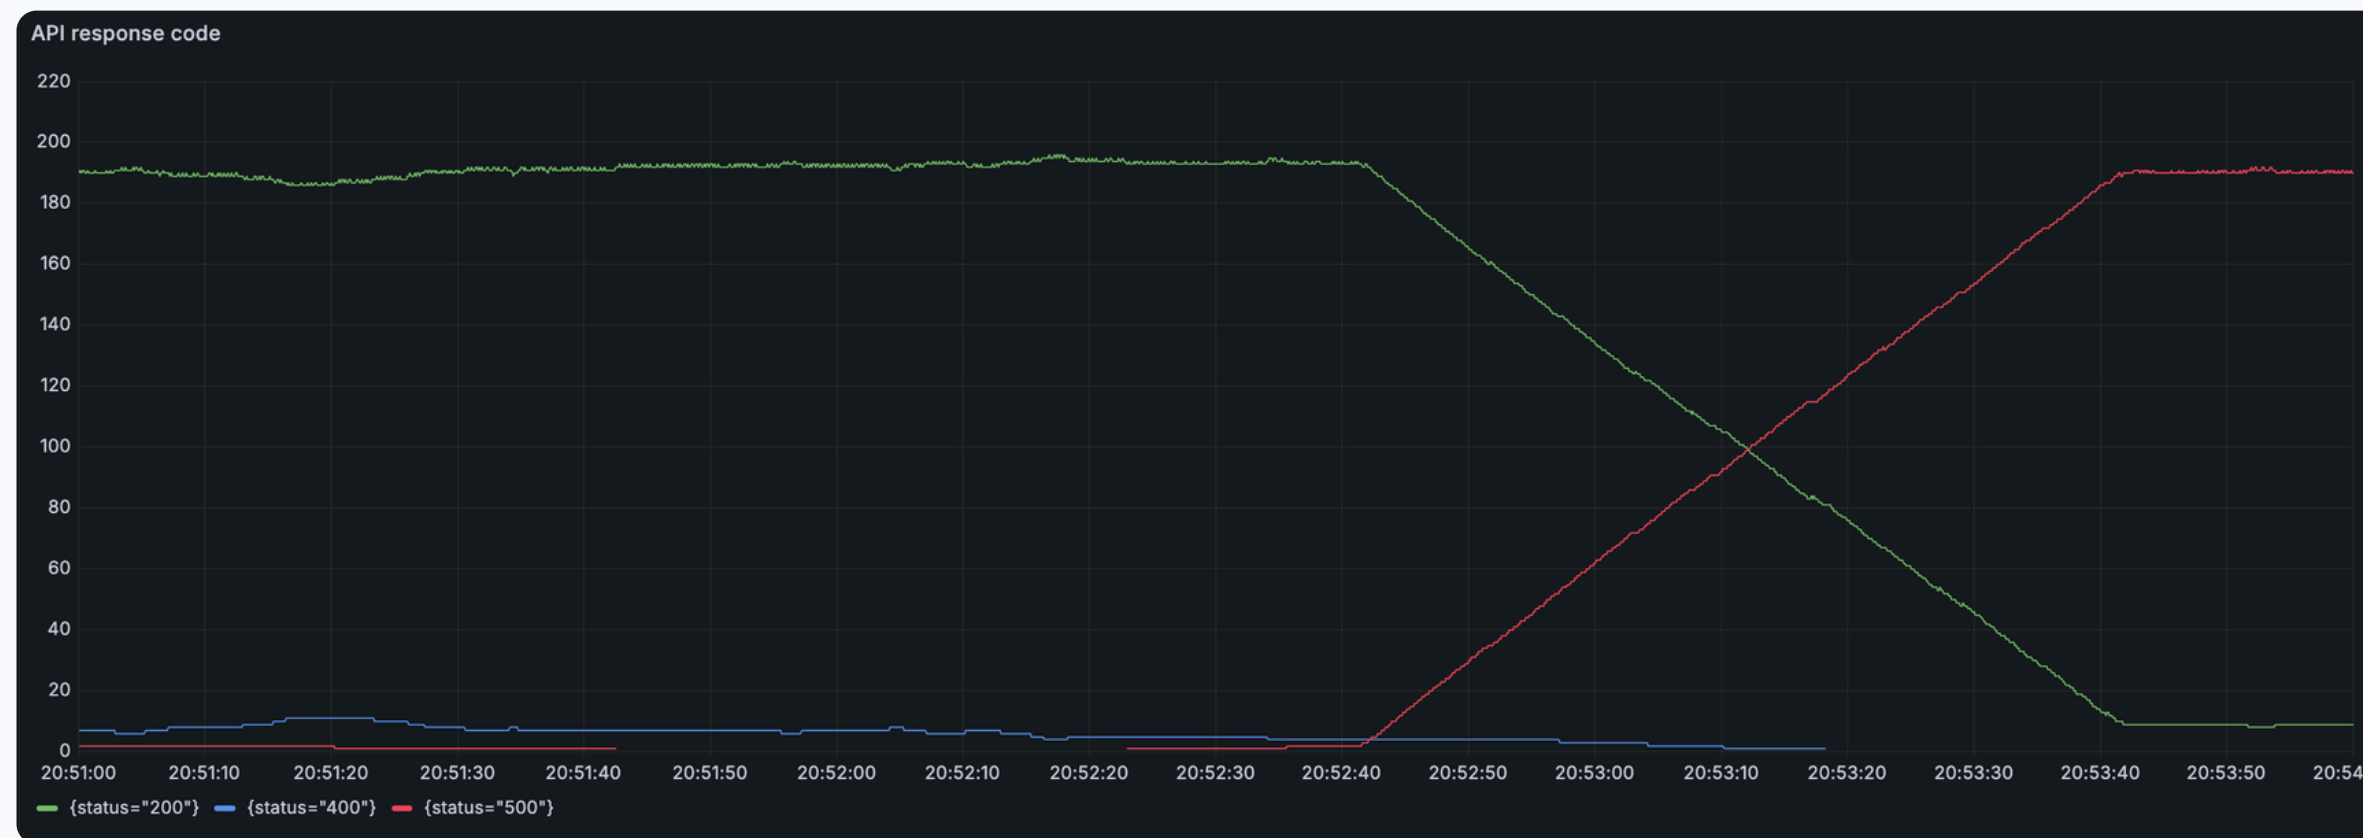
# How to Visualize Logs

1. Write the query
2. Apply transformations
3. Create variables
4. Add overrides



# Alerting Strategies

- Single error detection  
Trigger when a critical error appears in a log line.
- Error rate spikes (RED method)  
Alert when the rate of errors suddenly increases, signaling a systemic issue.



# Results

- Strong team adoption and enthusiasm.
- Significantly improved **MTTD** and reduced **MTTR** from 2 days to 2 hours, enhancing service quality.
- Handled 1 GB/day with 2-week retention (14 GB total), running smoothly on a €6/month instance.

# Thank you!

## Questions?

# Demo Repository

